

Board and C-level presentation: the secure agentic platform

Board and C-level presentation, 2026-10-01


What this asks of you

- Appoint the people; each tier of agents opens only when they exist
- Fund purchases nobody has quoted yet; start the long-lead ones now
- Make three signatures possible: the owner's two lists and purpose; a reviewer's deviation
- Accept eight residual risks in writing; row one is a tenant compromise
- Choose a start: three days, a proof of value, or the full build
- Nothing is built; no price is quoted; no regulator's decision is promised

Why now

- Objective of 2026-09-13: one secure environment; Wall-E does, Eve controls, Mo improves
- Wall-E holds Super Admin on a dedicated licensed user account (P33, decided)
- Hundreds of agents, future AGI-class agents, EU AI Act, TISAX
- Without a platform, each agent rebuilds its own security; nobody reviews it twice
- EU AI Act Articles 4, 5 and 50 bind already; penalties are real now
- The design declines "bulletproof" and AGI hosting in 2026; qualifies "any action"

What is proposed, in one picture



The central fact: nothing narrows a super admin

- Fact one: a service account cannot hold Super Admin; Wall-E is a user account
- Fact two: Super Admin cannot be limited to a unit or a privilege subset
- One enforcement point remains, the action service; Google's ceiling is the consented scopes
- A leaked token or interactive login is a tenant compromise reaching into Google Cloud
- Detection becomes the primary control; least privilege becomes a signed TISAX deviation
- The platform does not narrow a super admin; nothing in Workspace or Cloud can

What makes the grant acceptable

- Thirteen compensations are preconditions of the grant, none deferred; four checklist rows more
- Three lanes: only the catalogue runs unattended; uncatalogued work needs two human super admins
- The model never holds a credential; two credentials in two services; no domain-wide delegation
- No super-admin-class operation is ever autonomous; a human approves each one
- The grant is a gate: it happens when the last checklist row is green
- Two rows red, or detection, hygiene or second-human row red: Super Admin removed

Eve, stated honestly

- Eve's control path holds no model, by construction; it halts over plain REST
- Eve holds a read-only role, never Super Admin; it watches every human super admin
- Inside the organisation a super admin can reach Eve: independence there is detective
- Outside, the witness organisation holds evidence, incident record and pager: that is structural
- The witness sees Eve's deletion within fifteen minutes, not a forged approval
- Eve's control path relocates to the witness later (P15); not a grant condition

| How it can stop

- End of Stage 1: a dated stop-or-continue review; a negative balance stops the programme
- L3, a human approving each action, is a legitimate permanent end state
- Kill switches K0 to K7: halt writes, remove Super Admin, stop the fleet
- Silence is a halt: a quiet log pipeline halts the super-admin lanes
- The proof of value ends at a dated stop-or-continue record; nothing torn down
- The three-day build unwinds the same day, signed

The money

Cost	Driver	Paid by
Security Command Center Premium, from Tier C	organisation-wide findings service	IT security (Assumption: a platform line until then)
Central log ingestion and locked retention, from Tier R	Workspace audit volume, the largest line	the platform
Google SecOps with managed detection, from Tier P	events per day, retention, retainer	IT security
The witness organisation, from Tier P	one tenant, project, bucket, dataset, channels, two keys	IT security
Licences, hardware keys, Chrome Enterprise Premium (one licence per operator and approver)	one Workspace licence per Tier P agent; 22 or 26 keys	agent owner; the platform; procurement
Human hours, every tier	the roles on slide 9	the binding cost

The people you must appoint

Tier or gate	Distinct humans	Who is added
Tier C and Tier R	1	nobody new; self-review recorded as such
Tier W, the first write	3	second operator (about 2 h/week); security reviewer from IT security (about 2 h/month); blind grader (about 1 h/week per agent)
The super-admin grant	5 named, or 4 with a dated ISMS exception	second human super admin outside the Wall-E line, as Eve's owner; incident commander; engaged DPO; AI compliance owner with a deputy
Stage 3, autonomous catalogue work	4 plus a bought 24x7 desk	the detection desk that acknowledges within target
Steady-state super admins	exactly 3	two human admin accounts and walle@; the robot never the recovery one

Employees, the DPO and the works council

- Eve monitors the human super admins from its first run: that is employee monitoring
- A DPO record and HR's works-council answer come before Eve's first stream
- Stage 0 reads every employee's directory data; the DPIA precedes it, the DPO deciding
- Suspension only after HR or a named operator decided; restoring stays human-approved
- Jurisdiction unknown (`Assumption:`); consultation or consent may be required, not information
- The works-council question is the longest lead item; it goes out on day one

| The EU AI Act

- Classification follows the declared purpose, the catalogue, not the account's privilege
- Wall-E: adjacent to Annex III 4(b); narrow-task derogation claimed; registered before its first write
- Articles 4, 5 and 50 bind now; Annex III high-risk obligations from 2027-12-02
- Classification guidelines still draft on 2026-09-18; the assessment re-runs within 90 days of final
- Super Admin is the fact most likely held against the narrow reading
- Not promised: a regulator's reading; no harmonised standard, so no presumption of conformity

TISAX

- Target Confidential at AL2, remote evidence review with interview (Assumption:); a site earns it
- The robot's Super Admin deviates from control 4.2.1; a signed record with three signatories
- ISA2027 governs assessments ordered from 2027-01-01; ISA 6 orders end 2026-12-31
- One person in every role is the finding most likely to stop an assessment
- Nothing built in 2026 reaches maturity 3; the ISMS decides the cut-off by 2026-12-01
- Not promised: an assessor may refuse maturity 3 on least privilege despite the compensations

Three ways to start

Path	People and time	What it proves
The three-day build	two people, three business days, buys nothing	Eve live before any doer; forced dry run; two kill levers timed
The proof of value	three hands-on people, one engineer; 6 to 9 weeks to POV-1 (Assumption:)	Eve on every human super admin; a doer holding no admin role
The full build	thirteen appointments; 85 to 90 person-days; Stage 0 not before 2027-03	the platform; Eve with a witness; the gated grant; Stage 0

The timeline, by gate

When (Assumption: after 2026-09-24)	Gate or milestone	What must exist
2026-09-21, day one (not started on 2026-10-01)	the DPO letter, the works-council question, the toil baseline, every purchase started	the platform owner alone
2026-09-29 to 2026-10-27	decisions and people signed	the appointments of slide 9
2026-11-10 to 2027-01-19	Tier R open, then Tier C	Security Command Center Premium; nobody new
2026-12-22 to 2027-01-19	Eve's observe-and-report layer live over the human super admins	the DPO record; the witness; Eve's code
not before 2027-03	the super-admin grant, then Stage 0 one to two weeks later	the SIEM and MDR contract; the penetration test; five named humans; the checklist green
2027-12-02	Annex III obligations bind; Wall-E's later stages straddle it	Stage 5 at least 27 to 38 weeks after Stage 0

The risks you accept by signing

- A stolen robot credential is a tenant takeover; detection meets it, nothing prevents it
- Inside the organisation the controller watches a super admin but cannot structurally stop one
- The safety of the design depends on people who do not exist yet
- The robot's gate is one process; its build pipeline could rewrite it
- Manipulated content will reach the model; screening only detects; ceilings in code enforce
- The regulator and the assessor decide outcomes; some Google facts are unverified

What you sign first

- The gate-counting rule: how an open decision with a working fallback counts
- P29, the two lists: what is refused everywhere, what needs two humans (the owner)
- P28, Wall-E's declared intended purpose, signed with legal
- P33 and P136, the deviation record: security reviewer signs, the ISMS enters it
- P68, the super-admin roster of exactly three, with the second human named
- P34, Eve's reporting path may reason: the Eve owner and the security reviewer

The ask, restated

- Appoint: a second operator, a security reviewer, a blind grader, then the grant roles
- Fund: quotes for the monitoring service, the witness, the penetration test, the keys
- Enable: the owner signs lists and purpose; a reviewer and ISMS sign the deviation
- Accept: the eight residual risks, in writing, row one pending a security reviewer
- Choose: three days, the proof of value, or the full build; every date runs from that choice
- Send on day one: the DPO letter and the works-council question

| Appendix

The thirteen compensations, and the four rows the checklist adds

Rows	Precondition of the grant, none deferred
1 and 2	three bands in code; the two lists signed (P29), hash recorded in the deviation file
3 and 4	two credentials in two services, <code>cloud-platform</code> never consented; the band-B requester rule
5	detection as the primary control: Eve live and drilled, the witness receiving, a SIEM with 24x7 acknowledgement
6 and 7	account hygiene with Workspace multi-party approval on; kill switches drilled, K6 rehearsed on the sandbox twin
8, 9 and 10	the perimeter with PAM on the deploy grant; a permanent ceiling; the lost role scoping replaced in code
11, 12 and 13	a second human outside the Wall-E line; the signed record P33; the EU AI Act position P28

The kill switches

Switch	What it does	Who pulls it, and how fast
K0	halts writes	any operator, Eve or a breaker; containment under 60 seconds
K1 and K2	demotes one cell; stops new runs	K1 any operator, Eve or a breaker, as K0; K2 an operator with the cloud rights
K3 and K4	cuts the agent off; revokes the credential at Google	K3 a project administrator; K4 any operator, one call to each service; a token already issued may work up to 60 minutes (Assumption:)
K5	suspends the robot account or revokes the grant	a human super admin on the two-person rota, within 30 minutes of acknowledgement
K6	removes Super Admin from the robot	a human only, within 60 minutes; the switch that survives a minted token
K7	stops every agent at once	a human with the entitlement or a deterministic job on a severity-1 rule; never a model; drilled monthly

The register of decisions in one screen

Rows	Count	State on 2026-10-01
P1 to P34, the high-level design's own	34	3 decided or closed (P33 decided); 9 proposed; 18 open; 4 spikes
P35 to P143, the detailed pages'	109	107 proposed, 2 spikes; 11 carry an open dependency or an assumption
P144 to P191, the setup procedures'	48	all proposed, pending the owner's signature
P192 to P204, the proof of value's	13	all proposed; 6 POV-only, 6 binding on the full build, P204 mixed
P205 to P207, added on 2026-10-01	3	all proposed
Total	207	one row decided by the owner
Keeping the grant red	—	open P7, P10, P14, P29; P33's deviation signatures

The three build paths, compared

	The three-day build	The proof of value	The full build
People	two, both super admins, plus a sponsor	three hands-on, one engineer, named approvers	thirteen appointments
Time, effort, purchases	three business days, six person-days; buys nothing	60 to 90 person- and engineer-days; POV-1 ends 6 to 9 weeks after day one, POV-2 ends week 16 to 20 with a dedicated engineer; nothing bought on the critical path	85 to 90 person-days; 6 to 7 months to Stage 0; thirteen purchase rows
The doer's privilege	one scoped privilege over four synthetic accounts	none; no admin role at all	Super Admin, under the gate
Eve	live before the doer; inside the administrators' reach	reports on every human super admin; no witness	with the witness; independent proof by the second human
Exit	same-day unwind, signed	dated stop-or-continue record	Stage 1 stop-or-continue review; L3 as an end state
Can never claim	super-admin safety; Eve's independence; Model Armor evidence; minutes saved	super-admin safety; Eve's independence; that Model Armor blocked anything; minutes saved; that the doer does admin work; a tamper-proof audit	a regulator's or assessor's outcome

| The sentences never to be used

From the proof of value, verbatim, with its reasons: "1. 'Wall-E is safe as a super admin.' Nothing in Track A touches super-admin containment (G10, G11, G14, G20, Eve G-7 need Track B, §3). Nor may any report say or imply that the doer 'does admin work': it holds no Workspace admin role. 2. 'Eve is independent.' Independence in the design is structural (a witness organisation, a second tenant). The POV has neither (PV-D-03), and its Eve lives inside the reach of the administrators it watches. 3. 'Model Armor blocked the injection.' A probabilistic content screen is never a trust boundary, whatever its grade (../01-hld.md §0.2). It produced evidence." (pov/README.md §1.2)

From the three-day build, verbatim, with its reasons: "1. 'Wall-E is safe as a super admin.' Nothing here touches super-admin containment. No agent holds Super Admin at any moment of the three days. 2. 'Eve is independent.' Independence in the design is structural: a witness organisation, a second tenant. This set has neither (C-01), and its Eve lives inside the reach of the administrators it watches. 3. 'Model Armor blocked the injection.' A probabilistic content screen is never a trust boundary. And nothing here exercises it: no step sends anything through the floor, so these three days produce **no Model Armor evidence at all**, only the floor settings themselves (C-22). 4. 'We saved <n> minutes of admin work.' The doer acted only on synthetic accounts. No human minute was saved, and Mo's scorecard says so in bold (C-07)." (3-day/README.md §1.1)

For the full build and this deck, the phrase used in place of "bulletproof": the platform promises the mechanisms and the evidence, and lists what the regulator or the guidelines still decide (10-eu-ai-act.md §6).