

Sales presentation — the secure agentic platform

Sales presentation, 2026-09-18

■ The proposal in one sentence

- A platform on Google Cloud and Google Workspace for hundreds of agents.
- Three first tenants: Wall-E does, Eve controls, Mo improves.
- Offered: a design, three build procedures, a decision register; no product, no service.
- No price is quoted anywhere; the drivers and the payers are named.
- Giving the doer Super Admin is the designing organisation's decision, designed around.
- The ask is on slide 17: take the approach, count, decide, measure, price.

The problem: toil, agents arriving anyway, doing nothing

- Where the design was written: one administrator holds every role the platform needs.
- **Assumption:** weekly digests are made by hand today, or not at all.
- The licences already exist; no-code agents can be built inside Gemini Enterprise.
- No policy forces registration; unregistered agents are detected, and unreachable through the front door.
- Without a platform, the second agent re-implements the first agent's controls.
- Three EU AI Act articles already bind; penalty exposure is real now.

■ The promise, as the design states it

- Routine administration done faster and on the record.
- From the read-only stage: weekly digests of changes, unused accounts, unowned groups.
- From the first writing stage: leaver and joiner actions from chat; licence reclaim.
- At batch approval: one operator approves a whole planned run.
- A hundredth agent costs a register row and a factory run.
- No benefit figure exists yet; four weeks of measured toil produce it.

The three agents, and what each never does

Agent	What it does	What it never does
Wall-E, the doer	Catalogued Workspace administration a named operator asked for, through an account holding Super Admin	Decides who leaves or is promoted; selects a person by behaviour
Eve, the controller	Approves, halts and demotes over plain REST with no model on that path	Holds Super Admin; raises a level; writes to Workspace
Mo, the improver	Measures every agent from audit data and proposes changes by pull request	Holds a credential; reaches production except through a merged pull request

Why this design: enforced versus detected

- Every control is enforcement-grade (Google or code refuses) or detection-grade (something sees it).
- Nothing detection-grade stands alone in a safety argument.
- A probabilistic content screen is evidence, never a trust boundary.
- Nothing in Workspace or Google Cloud narrows a super-admin account.
- So for the super-admin tier, detection is the primary control, and stated.
- A leaked robot token is a tenant compromise; custody and detection bound it.


Why this design: a tier opens only when its people exist

Tier	What its agents may do	Adds
C, classical	No-code assistants in Gemini Enterprise; retrieval only; user's own permissions	Nobody new; Security Command Center Premium, organisation-wide
R, read tools	Read-only tools in the agent's own project	Nobody new; nothing more bought
W, write agents	Writes to a system of record through an action service	Second operator, part-time security reviewer, blind grader
P, privileged; P-SA, the singleton	A tenant-level admin credential; for P-SA, Super Admin	Second human super admin, incident commander, DPO; SIEM, witness, keys
X, AGI-class	Capabilities beyond the designers' foresight; not open	An AI-safety reviewer who does not exist


Why this design: the witness, and the second human

- A Workspace super admin can reach everything inside the organisation, Eve included.
- Inside the organisation Eve's independence is therefore detective, not structural.
- Evidence, incident record and pager go to a witness organisation outside the tenant.
- The witness sees Eve's deletion within fifteen minutes; it cannot see a forged approval.
- A second human outside the Wall-E line owns Eve and receives every page.
- Eve watches every human super admin from its first run, before Wall-E exists.

Why this design: humans raise autonomy, machines lower it



Proof on a path: three days, then weeks, then months



What the small paths cannot show

- Three days: no super admin for any agent, no witness, no SIEM.
- Three days produce no Model Armor evidence and save no human minute.
- The proof of value's doer holds no Workspace admin role at all.
- Neither path claims any line of the super-admin gate; only the full build does.
- A clause whose record is missing is struck from the claim, not softened.
- Four sentences are banned from every report and slide; they close this deck.

What it costs: drivers and payers, no price yet

Cost line	Driver	Paid by	From tier
Security Command Center Premium	Organisation-wide, every project	IT security (Assumption: a platform budget line until then)	C
Central log ingestion and locked retention, the largest line	Workspace audit volume plus Data Access logs	Platform	R
Google SecOps with a managed detection and response retainer	Events per day, retention months, the retainer	IT security	P
The witness organisation	One tenant, project, bucket, dataset, channels, two hardware keys	IT security	P
Licences and screening: one Workspace licence per Tier P agent; Chrome Enterprise Premium per operator; Model Armor per request	Seats, operators, requests	Agent owner; platform; agent budget	P; W; R
Human hours, the binding cost	The roles of slide 12	Every tier	Every tier

The people it needs, and who existed where it was designed

- Where it was designed, on 2026-09-13: one administrator; every role was one person.
- That is the finding most likely to stop a TISAX assessment.
- Distinct humans: one at Tier R, three at Tier W, four at the grant.
- At the grant, five named humans; four only with a dated ISMS exception.
- Tier W adds a second operator, a security reviewer and a blind grader, part-time.
- Tier P adds a second super admin, an incident commander, a DPO, a desk.

How it can stop

- A dated stop-or-continue review at the end of Stage 1, against operating cost.
- A human approving every action (L3) is a legitimate permanent end state.
- The proof of value ends in a dated review: continue, hold, or stop dormant.
- The three-day build unwinds the same day; nothing real is touched.
- Eight kill switches, from halting writes in under a minute to a fleet kill.
- After the grant, named compensation rows turning red pull K6 until green.

| The EU AI Act

- Classification follows the declared purpose, the catalogue, not the account's privilege.
- Wall-E: adjacent to Annex III 4(b); the narrow-task derogation claimed and documented.
- Registered in the EU database before the first write; no registration, no production.
- Art. 4, 5 and 50 bind now; Annex III obligations from 2027-12-02.
- Not promised: draft guidelines; no harmonised standard; Super Admin held against it.
- Every date re-verified 2026-09-18; the registration duty survived the Omnibus.

TISAX

- A label is earned by a site, never a platform; target Confidential, AL2 (Assumption:).
- A machine account holding Super Admin deviates from control 4.2.1, least privilege.
- A signed deviation: the owner decides, a reviewer signs, the ISMS enters it.
- Separation of duties: a counted minimum the gate enforces; one person today.
- ISA2027 governs orders from 2027-01-01; the last ISA 6 order date is 2026-12-31.
- Not promised: an assessor may refuse maturity 3 on 4.2.1 however good the compensation.

The residual risks a signature accepts

- A stolen robot credential is a tenant takeover; nothing prevents it, detection bounds it.
- Inside the organisation the controller watches a super admin but cannot stop one.
- The safety of the design depends on people who do not exist yet.
- The robot's gate is one process its build pipeline can rewrite.
- Manipulated content will reach the model; ceilings in code stop the action.
- The regulator and the assessor decide outcomes; some Google facts stay unverified.

The ask

- Take the approach, not a product: five things, one gate, two grades, banned sentences.
- Count your own roster first; one person in every role keeps the grant closed.
- Decide Super Admin for your doer yourselves; slide 5 follows only if you do.
- Start the toil baseline before any doer's first phase; no "before" exists afterwards.
- Obtain your own prices; none exist anywhere in this design.
- Send your works-council and DPO questions on day one.

| Appendix

What changes for another organisation

- The Super Admin decision (P33) is the designing organisation's; another decides its own.
- The starting roster (one administrator, no SOC) is the designing organisation's; re-count yours.
- Jurisdiction, works-council law, SOC and ISMS are unknown here; establish yours as facts.
- Existing licences, the Gemini Enterprise project and the Workspace edition change Tier C.
- Google's mechanisms; the nine promises and twelve demands port; two super-admin facts do not.
- What does not change: the constraints, the grades, the gate, the banned sentences.

The numbers behind the deck

Number	Value	Defined at
Decisions in the register	204 (P1–P204); one decided, P33; the rest proposed, open or spikes	the register in one screen
Distinct humans per stage	1 at Tier R, 3 at Tier W, 4 at the grant, 4 plus a bought desk at Stage 3; five is the clean state	the minimum per stage
Compensations at the grant	13, plus 4 checklist rows: penetration test, DPIA, tabletop, witnessed key custody	the checklist
Kill switches	8, K0 to K7; K7 has four levers	kill switches
Stage floors	S0 3–4 weeks, S1 4–6, S2 6–8, S3 6–8, S4 8–12; S5 at least 27–38 weeks after S0	stage overview
Hardware keys	22, or 26 with the twin robots; 18 firm	the key count

What the platform will not do

- It does not narrow a super-admin account; nothing in Google can.
- No model produces an approval, a signature, a halt, a veto or a refusal.
- It does not automate the Admin console under any robot session.
- It does not run a SOC; it buys acknowledgement and feeds an existing one.
- It grants domain-wide delegation to nothing, ever; it trains no model.
- It does not promise EU AI Act or TISAX outcomes; mechanisms and evidence only.

What each tier waits on

Tier	Opens when
C, classical	The register and the Gemini Enterprise baseline exist
R, read tools	The factory, the folder baseline, central logging and the shared registry exist
W, write agents	The autonomy contract, a platform verifier, a validator custodian, a nonprod folder, one restore drill; works-council information for any agent acting on employee accounts
P, privileged; P-SA, the singleton	Everything in W plus every precondition of the super-admin grant
X, AGI-class	Not open in 2026; a sandbox tier in the EU, a capability-evaluation report per model pin, an AI-safety reviewer, six months of fleet-kill drills, an independent-model monitor

The sentences never to be used

Copied from the two build paths that forbid them. They are listed here so that no slide, note, script or report about this platform uses them, and each carries the reason.

From the proof of value (never to be used):

1. "Wall-E is safe as a super admin." Nothing in Track A touches super-admin containment (G10, G11, G14, G20, Eve G-7 need Track B, §3). Nor may any report say or imply that the doer "does admin work": it holds no Workspace admin role.
2. "Eve is independent." Independence in the design is structural (a witness organisation, a second tenant). The POV has neither (PV-D-03), and its Eve lives inside the reach of the administrators it watches.
3. "Model Armor blocked the injection." A probabilistic content screen is never a trust boundary, whatever its grade (01-hld.md §0.2). It produced evidence.

From the three-day build (never to be used), the same three and one more:

1. "Wall-E is safe as a super admin." Nothing here touches super-admin containment. No agent holds Super Admin at any moment of the three days.
2. "Eve is independent." Independence in the design is structural: a witness organisation, a second tenant. This set has neither (C-01), and its Eve lives inside the reach of the administrators it watches.